



ENDGAME.

Sense. Discover. Act.TM

IDA Splode

Dynamically Enhancing Static Analysis



IDA Splode - Overview

ENDGAME.

- Background
- Motivation
- Walkthrough
- Examples
- Future
- Demo?

Background



IDA-Splode

- Horrible name for the tool I wrote, which combines...
- Hex-Rays IDA Pro
 - Application to disassemble software
 - **Static** analysis
 - `while(1){} → 0xebfe: jmp $-2`
- IDA Python
 - Python bindings for IDA's API
 - `while ea != BADADDR: ea = NextHead(ea)`
- Intel's Pin
 - Dynamic binary instrumentation framework
 - Effectively hook executing code at the instruction level
 - `INS_AddInstrumentFunction(CalledOnEachInstruction)`



Pin and Dynamic Binary Instrumentation

Dynamic Binary Instrumentation (DBI) is a method of analyzing the behavior of a binary application at runtime through the injection of instrumentation code. This instrumentation code executes as part of the normal instruction stream after being injected.

– Skape (Uninformed v7a1)



ENDGAME.

Motivation



Origins of IDA Splode

- Single-byte mutation
 - Heap overflow with corrupted data
 - Decompression failure
- Root Cause Analysis of Heap Overflow
 - Parser for undocumented format
 - Closed source, but have debugging symbols
 - All file data is compressed (flat entropy)
 - Extensive use of COM interfaces
 - Lots of C++ isms (OO, virtual inheritance, etc)
 - Lots of shiny things to get distracted by
- Too much churning, need better tools



Motivation for Developing IDA Splode

- IDAPython
 - Limited exposure
 - Small-scale applications
- Pin
 - Powerful tool with lots of potential
 - No previous personal experience
 - Some previous organizational experience
- Exists to answer
 - Given `mov [esi+0x18], eax`
 - What is in the memory `esi` points to?
 - Where did the allocation come from?
 - Which routines touch that data?



Static Reversing & Runtime Analysis

- Load binary in Windbg
 - Breakpoint on instruction
 - !heap, kv, ln, dd, wt
 - All require useful symbols, standard allocators
- Load binary in IDA
 - Annotate nearby instances of `esi`
 - Maybe create a struct
 - Tedious, manual, error-prone
- Single-step sample size $n=1$, no automation



Dynamically Enhanced Analysis

- PIN Tool + IDAPython
- Log all memory accesses
 - Instrument all `mov/movzx/lea/cmp`
- Record metadata at runtime
 - Stack?
 - Instrument all `call` and `ret` and instructions
 - Given stack pointer, know parent frame & BP offset
 - Heap?
 - Page Heap gives allocation size and stack trace
 - Given heap pointer, know size, offset and origin
 - Symbol?
 - Windbg/DbgHelp.dll to resolve symbols



Scope and Features

- Scope Limited by Pin's OS Support
 - Intel CPUs only (Windows, Linux, OSX, Android-on-Intel)
 - Currently Windows-only
 - For other OSes, would need to change
 - Debug symbol resolution (macho, elf)
 - Heap allocation tagging
- Does
 - Augment reversing (no source)
 - Make getting started easier
 - Get more useful with debug symbols
- Does Not
 - Detect OOB access or exploitable conditions (a la Valgrind)
 - Track taints, symbolic, concolic, etc. analysis
 - If you see “symbol” or “symbolic” in this presentation, think Windbg



Basic Use Walkthrough



Three Parts to IDA Splode

- Running Pin Instrumentation
- Import Traces Into Database
- IDA Pro Script, Usage, and Hotkeys



Running Pin Instrumentation

- Acquire `ida-splode.dll` pintool (on SVN)
- Run target binary under Pin, specify `ida-splode.dll` tool
 - `pin.exe -t ida-splode.dll -- target.exe args`
- Outside scope of presentation, knobs available for...
 - Metadata gathering
 - Heap vs. Stack vs. Debug Symbols
 - Limiting scope of instrumentation or logging
 - By module (`-m foo.dll`) or routine (`-r foo!Function`)
 - Types of instructions instrumented
 - MOV vs. LEA vs. CMP
 - Instrument registers vs. const values



Import Traces into Database

- Uses MongoDB for ease-of-use and zero configuration
 - `\path\to\bin\mongod --dbpath "any folder"`
- Log files created by Pin instrumentation
 - All logs are plain Python
 - One dict per instrumented instruction
 - Greppable
 - Run `target.exe.py` to import traces into MongoDB
 - Traces are in `target.exe.py.traces.py`
 - Modules are in `target.exe.py.modules.py`



IDA Pro Script & Usage

- Single script to run via ALT+F7 or “File>Script File”
 - `ida-splode\py\idapython_script.py`
- Upon execution
 - Displays usage and hotkeys
 - Connects to MongoDB on localhost
 - Shows a list of candidate databases
 - Must contain a module with matching MD5
 - First run, asks user which database to load
 - Subsequent runs, auto-loads previously-used database
- Common Hotkeys
 - ^↑H Help and usage
 - ^↑C Color instructions in module
 - ^↑S Summary for selected instruction
 - ^↑D Details for selected instruction
 - ^↑X Create X-refs from module trace info
 - ^↑L Load a different MongoDB database



Examples & Screen Shots



Example Usage Cases

- Easiest to demonstrate in bite-sized, contrived examples
- Features Demonstrated
 - Instruction Tracing
 - Memory Value Recording
 - Branch Tracing & Statistics
 - Vtable Resolution
 - IAT Rebuilding
 - Extramodular xrefs
 - Branch Statistics
 - Stack Variable Propagation
 - Heap Metadata



IDA Splode - Examples

ENDGAME.

Instruction Tracing

- Simplest PIN task
 - On each instruction, log IP
- Simplest IDA task
 - Color those instructions

Example Disassembly

```
push    ebp
mov     ebp, esp
sub     esp, 0Ch
push    4 ; unsigned int
call    operator new(uint)
add     esp, 4
mov     [ebp+var_8], eax
cmp     [ebp+var_8], 0
jz      short loc_E510A6
```

```
mov     ecx, [ebp+var_8]
call    sub_E51520
mov     [ebp+var_C], eax
jmp     short loc_E510AD
```

```
loc_E510A6:
mov     [ebp+var_C], 0
```

```
loc_E510AD:
mov     eax, [ebp+var_C]
mov     [ebp+var_4], eax
mov     ecx, [ebp+var_4]
mov     edx, [ecx]
mov     ecx, [ebp+var_4]
mov     eax, [edx+4]
call    eax
mov     esp, ebp
pop     ebp
retn
```

Enhanced

Executed

Not
Executed



Memory Values

- For most instructions...
 - Log memory value read/set
 - Log memory address
- Basic stats on common vals
- Values stand out

Example Disassembly

- Normal Run

```
mov     [ebp+var_14], ecx ; >< 24h
```

- Crashing Run

```
mov     [ebp+var_14], ecx ; >< 8080A47Ch(50%), 24h(50%)
```



IDA Splode - Examples

ENDGAME.

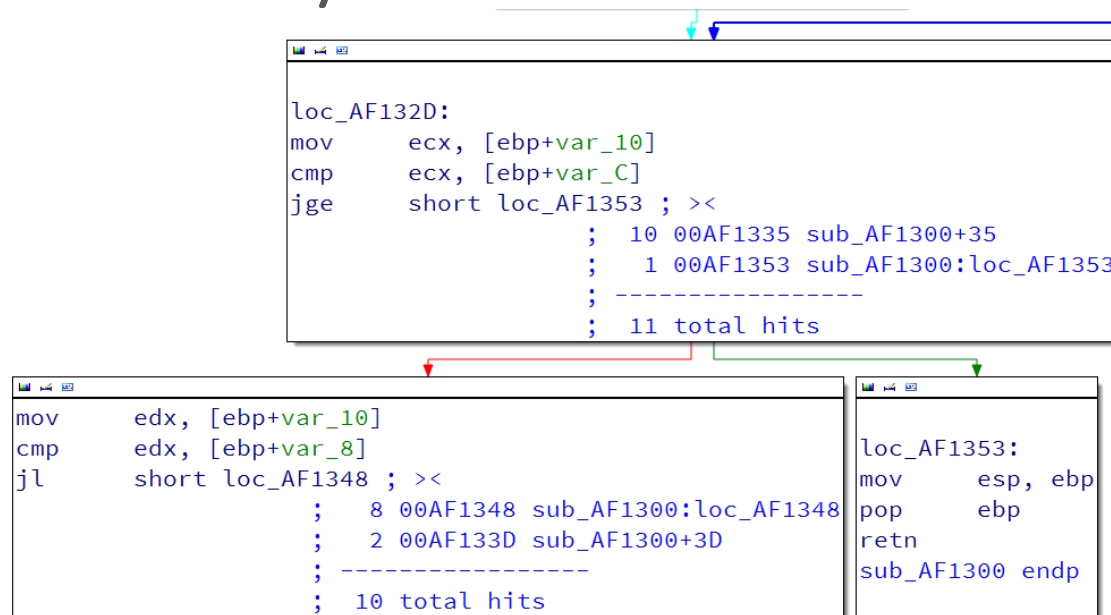
Branch Tracing

- Logs branches taken
- Generates statistics
- Useful for
 - Loops
 - Error handlers

Example Source

```
void TestBranchStatistics() {  
    int x = 0, ten=10, eight=8;  
    for(int i = 0; i < ten; i++)  
    {  
        if(i >= eight) { x++; }  
        else           { x--; }  
    }  
}
```

Example Disassembly





Vtable Call

- Lots of Pin Tools do this
 - Might as well do it too
- Adds XREFs

Example Source

```
struct V {  
    virtual ~V(){};  
    virtual void Foo()=0;  
};  
struct V1 : public V { void Foo(){} };  
void TestVtable() {  
    V *v = new V1;  
    v->Foo();  
}
```

Example Disassembly

```
mov     ecx, [ebp+v]  
mov     eax, [edx+4]  
call    eax ; >< 1 00E51540 V1::Foo
```



Dynamic Imports

- Lots of Pin Tools do this
 - Might as well do it too
- Dynamic Import Section
 - Adds function stub
 - Adds XREFs

Example Source

```
typedef void (WINAPI * PfnODS)(char*);  
PfnODS pODS = NULL;  
#define szMod      "kernel32.dll"  
#define szFnName   "OutputDebugStringA"  
void TestDynamicIAT() {  
    HMODULE h = GetModuleHandleA(szMod);  
    pODS = (PfnODS) GetProcAddress(h, szFnName);  
    pODS("0x1239");  
}
```

Example Disassembly

```
call    dword_B8345C    ; >< 1 00B94000 KERNEL32!OutputDebugStringA
```



Extramodular Xrefs

- Dynamic imports is only half the battle
- Where else is <API> called?
- Segment created for external refs
 - Create Fn stub for each enhanced insn
 - ...or batch mode (entire module)

Example Disassembly

```
.idasplode:00B94000 KERNEL32_OutputDebugStringA: ; CODE XREF: sub_B81290+2B↑p  
.idasplode:00B94000         retn     4
```




Stack Tracking

- Tracks all stack frames
 - Hook each *call/ret*
 - Save ESP after *call*
 - Check ESP after *ret*
- Ptr-to-stack = obvious

Example Disassembly

```
mov     eax, [ebp+C] ; >< 00D41260 TestStackVars A
mov     dword ptr [eax], 1237h
```

Example Source

```
void stack_var_user(int* C) {
    *C = 0x1237;
}
void stack_var_middleman(int* B) {
    stack_var_user(B);
}
void TestStackVars() {
    int A = 0x1238;
    stack_var_middleman(&A);
}
```



Heap Metadata

- Contrived Example
 - Create struct on heap
 - Set values inside struct
 - `__thiscall` and `__stdcall`

Example Source

```
#include <pshpack1.h>
struct S {
    char a;
    short b;
    int c;
    S() { c = 0x1230;
        b = 0x1231;
        a = 1; }
    void setA() { a = 1; }
    void setB() { b = 0x1232; }
    void setC() { c = 0x1233; }
};
#include <poppack.h>

void UseS(S* s) { s->c = 0x1234; }

void TestStruct() {
    S *s = new S();
    s->setA();
    s->setB();
    s->setC();

    UseS(s);

    delete s;
}
```



Heap Metadata

- var_8 looks like a object
 - Where is it from?
 - How big is it?
 - Where else is it used?
- How can IDA Splode help?

Example Disassembly

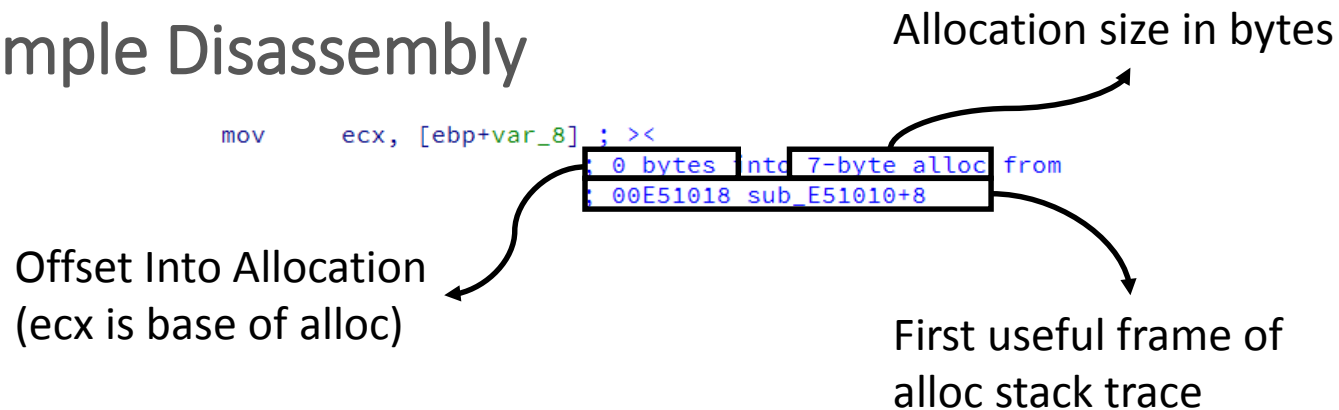
```
mov     ecx, [ebp+var_8]
call    sub_B81490
mov     [ebp+var_10], eax
jmp     short loc_B8103D
```



Summary View

- “Value” only
- Single hotkey to enhance
- Smaller amount of screen space

Example Disassembly





Details View

- “Address” & “Value”
 - Does what it says on tin
- Heap metadata
 - More exploded view
 - Statistics & Stack Trace

Example Disassembly

```
mov     ecx, [ebp+var_8] ; ><
; ===== Address =====
; => Stack Pointers
;      00E51010 sub_E51010 var_8
; ===== Value =====
; => Heap Pointers
; Count  Pctg  Offset Size
;      1  100%    0x0  0x7
; Backtrace:
;      00E51018 sub_E51010+8
;      00E5142B _main+1B
call    sub_E51490
```

Where are we reading
from or writing to?

What was read from or
written to that address?



Details View

- “Address” & “Value”
 - Does what it says on tin
- Heap metadata
 - More exploded view
 - Statistics & Stack Trace

Example Disassembly

```
mov     ecx, [ebp+var_8] ; ><
; ===== Address =====
; => Stack Pointers
;      00E51010 sub_E51010 var_8
; ===== Value =====
; => Heap Pointers
; Count  Pctg  Offset Size
;      1   100%    0x0  0x7
; Backtrace:
;      00E51018 sub_E51010+8
;      00E5142B _main+1B
call    sub_E51490
```

Tells us what we already know.

Heap!

One distinct allocation size/offset pair, which is seen 100% of the time.

Allocation stack trace; trims off actual allocator name (e.g. malloc/new). Addresses clickable.



Heap Metadata

- var_8 looks like an object
 - Created at sub_E51010+8
 - Easy search forward for c'tor
- How big is it?
 - 7 bytes
- What does it look like?
- Where else is it used?

```
mov     ecx, [ebp+var_8] ; ><
; 0 bytes into 7-byte alloc from
; 00E51018 sub_E51010+8
```



Summary Hotkey

```
mov     ecx, [ebp+var_8]
call    sub_E51490
```



Details Hotkey

```
mov     ecx, [ebp+var_8] ; ><
; ===== Address =====
; => Stack Pointers
;      00E51010 sub_E51010 var_8
; ===== Value =====
; => Heap Pointers
; Count  Pctg  Offset Size
;      1  100%    0x0  0x7
; Backtrace:
;      00E51018 sub_E51010+8
;      00E5142B _main+1B
call    sub_E51490
```



Heap Metadata

- var_8 looks like an object
 - Created at sub_E51010+8
 - Easy search forward for c'tor
- How big is it?
 - 7 bytes
- What does it look like?
 - Three fields (1b, 2b, 4b)
- Where else is it used?
 - sub_E51490
 - sub_E514C0
 - Etc.

```
struct S {  
    char a;  
    short b;  
    int c;  
};
```

↓ Compile

```
mov ecx, [ebp+var_8]  
call sub_E51490
```

↓ Reconstruct Hotkey

```
00000000 MyStruct struc ; (sizeof=0x7)  
00000000 autofield_0 db ? ; XREF: sub_E51490+20↑w sub_E514C0+A↑w  
00000001 autofield_1 dw ? ; XREF: sub_E51490+19↑w sub_E514E0+F↑w  
00000003 autofield_3 dd ? ; XREF: sub_E51000+6↑w sub_E51490+A↑w ...  
00000007 MyStruct ends
```




Vtable Call (Again)

- Same Example
- Heap metadata enhanced
 - Provides a bigger picture

Source

```
struct V {  
    virtual ~V(){};  
    virtual void Foo()=0;  
};  
struct V1 : public V { void Foo(){} };  
void TestVtable() {  
    V *v = new V1;  
    v->Foo();  
}
```

```
mov     ecx, [ebp+v] ; ><  
        ; 0 bytes into 4-byte alloc from  
        ; 00E51088 TestVtable+8  
mov     eax, [edx+4] ; ><  
        ; ===== Address =====  
        ; => Symbols  
        ;      00E5217C const V1::`vftable'+4  
        ; ===== Value =====  
        ; => Symbols  
        ;      00E51540 V1::Foo  
call    eax ; >< 1 00E51540 V1::Foo
```



Scenario: Custom Allocator

- Can't, or won't, change to PageHeap-compatible
 - E.g. Closed-source software implementing its own heap
- Redirect execution at runtime
- In this example, *custom_malloc* redirected to *malloc*

```
int heap[0x100];  
void* custom_malloc(int size) { return &heap[0x80]; }
```

```
call    custom_malloc(int)  
add     esp, 4  
mov     [ebp+memory], eax ; ><  
; ===== Address =====  
; => Stack Pointers  
;      00E511B0 TestCustomMalloc memory  
; ===== Value =====  
; => Heap Pointers  
; Count  Pctg  Offset Size  
;      1  100%    0x0 0x1236  
; Backtrace:  
;      00E511E8 TestCustomMalloc+38  
;      00E5144E _main+3E
```

Works because
of redirection
to *malloc*



Scenario: Custom Allocator

- Can't, or won't, change to PageHeap-compatible
 - E.g. Closed-source software implementing its own heap
- Redirect execution at runtime
- In this example, *custom_malloc* redirected to *malloc*

```
// -- Demo application allocator hooks  
new SymbolResolver<INS>(new MallocPageHeapRedirector(0,1), "demo.exe!custom_malloc");
```



Moveable Heap

- GlobalAlloc & LocalAlloc
 - Allow MEM_MOVEABLE flag
 - Used by IStream COM iface
 - E.g. lots of MSFT code
- Get handle to allocation
 - OS can move around allocation & data unless “Locked”
- No PageHeap Support
 - Hook RtlAllocateHandle
 - Keep track at runtime
 - HGLOBALs are heap pointers, not possible to confuse a la HANDLES

```
void TestMoveableHGLOBAL() {  
    HGLOBAL g = GlobalAlloc(GHND, 0x123a);  
    void * m = GlobalLock(g);  
    GlobalUnlock(m);  
    GlobalFree(g);  
}
```

```
push    123Ah    ; dwBytes  
push    GMEM_MOVEABLE or GMEM_ZEROINIT ; uFlags  
call    ds:GlobalAlloc  
mov     [ebp+hMem], eax ; ><  
                ; 0 bytes into 123Ah-byte alloc from  
                ; 00E512DD sub_E512D0+D
```



ENDGAME.

Future



Potential Enhancements & Future

- Scalability & Analytics
 - Lots of trace data generated every run
 - 80mb of data for ~0.030 seconds of computation
- Defense
 - CTFGRIND + EMET emulation
- Valgrind Rewrite
 - Useful for ARM binaries
- Linux/OS X Support
 - Not everything is Windows-based
- More intuitive/useful UI
 - Comments are quick-and-easy
 - IDA runs a full QT UI



ENDGAME.

Questions?



ENDGAME.

Sense. Discover. Act.TM